

# FedAvg算法复现

## 1. 准备工作

FedAvg算法过程如下:

---

**Algorithm 1** FederatedAveraging. The  $K$  clients are indexed by  $k$ ;  $B$  is the local minibatch size,  $E$  is the number of local epochs, and  $\eta$  is the learning rate.

---

**Server executes:**

initialize  $w_0$

**for** each round  $t = 1, 2, \dots$  **do**

$m \leftarrow \max(C \cdot K, 1)$

$S_t \leftarrow$  (random set of  $m$  clients)

**for** each client  $k \in S_t$  **in parallel do**

$w_{t+1}^k \leftarrow \text{ClientUpdate}(k, w_t)$

$w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$

**ClientUpdate( $k, w$ ):** // Run on client  $k$

$\mathcal{B} \leftarrow$  (split  $\mathcal{P}_k$  into batches of size  $B$ )

**for** each local epoch  $i$  from 1 to  $E$  **do**

**for** batch  $b \in \mathcal{B}$  **do**

$w \leftarrow w - \eta \nabla \ell(w; b)$

    return  $w$  to server

---

数据集介绍:

CIFAR-10是一个更接近普适物体的彩色图像数据集。CIFAR-10 是由Hinton 的学生Alex Krizhevsky 和 Ilya Sutskever 整理的一个用于识别普适物体的小型数据集。一共包含10 个类别的RGB 彩色图片: 飞机 (airplane)、汽车 (automobile)、鸟类 (bird)、猫 (cat)、鹿 (deer)、狗 (dog)、蛙类 (frog)、马 (horse)、船 (ship) 和卡车 (truck)。每个图片的尺寸为 $32 \times 32$ , 每个类别有6000个图像, 数据集中一共有50000 张训练图片和10000 张测试图片。

## 2. 分割数据集

```
1 def get_datasets(data_name, dataroot, normalize=True, val_size=10000):
2     """
3     get_datasets returns train/val/test data splits of CIFAR10/100 datasets
4     :param data_name: name of dataset, choose from [cifar10, cifar100]
5     :param dataroot: root to data dir
```

```

6      :param normalize: True/False to normalize the data
7      :param val_size: validation split size (in #samples)
8      :return: train_set, val_set, test_set (tuple of pytorch dataset/subset)
9      """
10
11     if data_name == 'cifar10':
12         normalization = transforms.Normalize((0.4914, 0.4822, 0.4465),
13 (0.2023, 0.1994, 0.2010))
14         data_obj = CIFAR10
15     elif data_name == 'cifar100':
16         normalization = transforms.Normalize((0.5071, 0.4865, 0.4409),
17 (0.2673, 0.2564, 0.2762))
18         data_obj = CIFAR100
19     else:
20         raise ValueError("choose data_name from ['mnist', 'cifar10',
21 'cifar100']")
22
23     trans = [transforms.ToTensor()]
24
25     if normalize:
26         trans.append(normalization)
27
28     transform = transforms.Compose(trans)
29
30     dataset = data_obj(
31         dataroot,
32         train=True,
33         download=True,
34         transform=transform
35     )
36
37     test_set = data_obj(
38         dataroot,
39         train=False,
40         download=True,
41         transform=transform
42     )
43
44     train_size = len(dataset) - val_size
45     train_set, val_set = torch.utils.data.random_split(dataset,
46 [train_size, val_size]) # 切割数据集为训练集与验证集
47
48     return train_set, val_set, test_set
49
50 def get_num_classes_samples(dataset):
51     """
52     extracts info about certain dataset
53     :param dataset: pytorch dataset object
54     :return: dataset info number of classes, number of samples, list of
55 labels
56     """
57     # -----#
58     # Extract labels #
59     # -----#
60     if isinstance(dataset, torch.utils.data.Subset):
61         if isinstance(dataset.dataset.targets, list):

```

```

58         data_labels_list = np.array(dataset.dataset.targets)
[dataset.indices]
59     else:
60         data_labels_list = dataset.dataset.targets[dataset.indices]
61     else:
62         if isinstance(dataset.targets, list):
63             data_labels_list = np.array(dataset.targets)
64         else:
65             data_labels_list = dataset.targets
66     classes, num_samples = np.unique(data_labels_list, return_counts=True)
67     num_classes = len(classes)
68     return num_classes, num_samples, data_labels_list
69
70
71 def gen_classes_per_node(dataset, num_users, classes_per_user=2,
high_prob=0.6, low_prob=0.4):
72     """
73     creates the data distribution of each client
74     :param dataset: pytorch dataset object
75     :param num_users: number of clients
76     :param classes_per_user: number of classes assigned to each client
77     :param high_prob: highest prob sampled
78     :param low_prob: lowest prob sampled
79     :return: dictionary mapping between classes and proportions, each entry
refers to other client
80     """
81     num_classes, num_samples, _ = get_num_classes_samples(dataset)
82
83     # -----#
84     # Divide classes + num samples for each user #
85     # -----#
86     assert (classes_per_user * num_users) % num_classes == 0, "equal
classes appearance is needed"
87     count_per_class = (classes_per_user * num_users) // num_classes
88     class_dict = {}
89     for i in range(num_classes):
90         # sampling alpha_i_c
91         probs = np.random.uniform(low_prob, high_prob,
size=count_per_class)
92         # normalizing
93         probs_norm = (probs / probs.sum()).tolist()
94         class_dict[i] = {'count': count_per_class, 'prob': probs_norm}
95
96     # -----#
97     # Assign each client with data indexes #
98     # -----#
99     class_partitions = defaultdict(list)
100     for i in range(num_users):
101         c = []
102         for _ in range(classes_per_user):
103             class_counts = [class_dict[i]['count'] for i in
range(num_classes)]
104             max_class_counts = np.where(np.array(class_counts) ==
max(class_counts))[0]
105             c.append(np.random.choice(max_class_counts))
106             class_dict[c[-1]]['count'] -= 1
107             class_partitions['class'].append(c)

```

```

108         class_partitions['prob'].append([class_dict[i]['prob'].pop() for i
in c])
109     return class_partitions
110
111
112 def gen_data_split(dataset, num_users, class_partitions):
113     """
114     divide data indexes for each client based on class_partition
115     :param dataset: pytorch dataset object (train/val/test)
116     :param num_users: number of clients
117     :param class_partitions: proportion of classes per client
118     :return: dictionary mapping client to its indexes
119     """
120     num_classes, num_samples, data_labels_list =
get_num_classes_samples(dataset)
121
122     # ----- #
123     # Create class index mapping #
124     # ----- #
125     data_class_idx = {i: np.where(data_labels_list == i)[0] for i in
range(num_classes)}
126
127     # ----- #
128     # Shuffling #
129     # ----- #
130     for data_idx in data_class_idx.values():
131         random.shuffle(data_idx)
132
133     # ----- #
134     # Assigning samples to each user #
135     # ----- #
136     user_data_idx = [[] for i in range(num_users)]
137     for usr_i in range(num_users):
138         for c, p in zip(class_partitions['class'][usr_i],
class_partitions['prob'][usr_i]):
139             end_idx = int(num_samples[c] * p)
140             user_data_idx[usr_i].extend(data_class_idx[c][:end_idx])
141             data_class_idx[c] = data_class_idx[c][end_idx:]
142
143     return user_data_idx
144
145
146 def gen_random_loaders(data_name, data_path, num_users, bz,
classes_per_user):
147     """
148     generates train/val/test loaders of each client
149     :param data_name: name of dataset, choose from [cifar10, cifar100]
150     :param data_path: root path for data dir
151     :param num_users: number of clients
152     :param bz: batch size
153     :param classes_per_user: number of classes assigned to each client
154     :return: train/val/test loaders of each client, list of pytorch
dataloaders
155     """
156     loader_params = {"batch_size": bz, "shuffle": False, "pin_memory":
True, "num_workers": 0}
157     dataloaders = []
158     datasets = get_datasets(data_name, data_path, normalize=True)

```

```

159     for i, d in enumerate(datasets):
160         # ensure same partition for train/test/val
161         if i == 0:
162             cls_partitions = gen_classes_per_node(d, num_users,
classes_per_user)
163             loader_params['shuffle'] = True
164             usr_subset_idx = gen_data_split(d, num_users, cls_partitions)
165             # create subsets for each client
166             subsets = list(map(lambda x: torch.utils.data.Subset(d, x),
usr_subset_idx))
167             # create dataloaders from subsets
168             dataloaders.append(list(map(lambda x:
torch.utils.data.DataLoader(x, **loader_params), subsets)))
169
170     return dataloaders
171

```

### 3. 数据节点类

```

1  from experiments.dataset import gen_random_loaders
2
3
4  class BaseNodes:
5      def __init__(
6          self,
7          data_name,
8          data_path,
9          n_nodes,
10         batch_size=128,
11         classes_per_node=2
12     ):
13
14         self.data_name = data_name
15         self.data_path = data_path
16         self.n_nodes = n_nodes
17         self.classes_per_node = classes_per_node
18
19         self.batch_size = batch_size
20
21         self.train_loaders, self.val_loaders, self.test_loaders = None,
None, None
22         self._init_dataloaders()
23
24     def _init_dataloaders(self):
25         self.train_loaders, self.val_loaders, self.test_loaders =
gen_random_loaders(
26             self.data_name,
27             self.data_path,
28             self.n_nodes,
29             self.batch_size,
30             self.classes_per_node
31         )
32
33     def __len__(self):
34         return self.n_nodes

```

## 4. CNN模型类

```
1 import torch.nn.functional as F
2 from torch import nn
3 import numpy as np
4 import torch
5 from torch.utils.data import TensorDataset
6 from torch.utils.data import DataLoader
7
8
9 class CNN(nn.Module):
10     def __init__(self, in_channels=3, n_kernels=16, out_dim=10):
11         super(CNN, self).__init__()
12
13         self.conv1 = nn.Conv2d(in_channels=in_channels,
14 out_channels=n_kernels, kernel_size=5)
15         self.pool = nn.MaxPool2d(2, 2)
16         self.conv2 = nn.Conv2d(in_channels=n_kernels, out_channels=2 *
17 n_kernels, kernel_size=5)
18         self.fc1 = nn.Linear(in_features=2 * n_kernels * 5 * 5,
19 out_features=120)
20         self.fc2 = nn.Linear(in_features=120, out_features=84)
21         self.fc3 = nn.Linear(in_features=84, out_features=out_dim)
22
23     def forward(self, x):
24         x = self.pool(F.relu(self.conv1(x)))
25         x = self.pool(F.relu(self.conv2(x)))
26         x = x.view(x.shape[0], -1)
27         x = F.relu(self.fc1(x))
28         x = F.relu(self.fc2(x))
29         x = self.fc3(x)
30         return x
31
32 class Client(object):
33     def __init__(self, trainDataSet, dev):
34         self.train_ds = trainDataSet
35         self.dev = dev
36         self.train_dl = None
37         self.local_parameter = None
```

## 5. 利用FedAvg算法训练

```
1 def train(data_name: str, data_path: str, classes_per_node: int, num_nodes:
2 int,
3         steps: int, node_iter: int, optim: str, lr: float, inner_lr:
4 float,
5         embed_lr: float, wd: float, inner_wd: float, embed_dim: int,
6 hyper_hid: int,
7         n_hidden: int, n_kernels: int, bs: int, device, eval_every: int,
8 save_path: Path,
9         seed: int) -> None:
10     #####
```

```

7      # init nodes, hnet, local net #
8      #####
9      steps = 5
10     node_iter = 5
11     nodes = BaseNodes(data_name, data_path, num_nodes,
12                        classes_per_node=classes_per_node,
13                        batch_size=bs)
14     net = CNN(n_kernels=n_kernels)
15     # hnet = hnet.to(device)
16     net = net.to(device)
17
18     #####
19     # init optimizer #
20     #####
21     # embed_lr = embed_lr if embed_lr is not None else lr
22     optimizer = torch.optim.SGD(
23         net.parameters(), lr=inner_lr, momentum=.9, weight_decay=inner_wd
24     )
25     criteria = torch.nn.CrossEntropyLoss()
26
27     #####
28     # init metrics #
29     #####
30     # step_iter = trange(steps)
31     step_iter = range(steps)
32     # train process
33     # record the global parameters
34     global_parameters = {}
35     for key, parameter in net.state_dict().items():
36         global_parameters[key] = parameter.clone()
37     for step in step_iter:
38
39         local_parameters_list = {}
40         # 需要训练的node数目
41         for i in range(node_iter):
42             # 随机选择一个客户端
43             node_id = random.choice(range(num_nodes))
44             # 用全局模型参数训练当前客户端
45             local_parameters = local_upload(nodes.train_loaders[node_id], 5,
46                                             net, criteria, optimizer,
47                                             global_parameters, dev='cpu')
48             print("\nEpoch: {}, Node Count: {}, Node ID: {}".format(step +
49                               1, i + 1, node_id), end="")
50             evaluate(net, local_parameters, nodes.val_loaders[node_id],
51                     'cpu')
52             local_parameters_list[i] = local_parameters
53
54         # 更新当前轮次模型的参数
55         sum_parameters = None
56         for node_id, parameters in local_parameters_list.items():
57             if sum_parameters is None:
58                 sum_parameters = parameters
59             else:
60                 for key in parameters.keys():
61                     sum_parameters[key] += parameters[key]
62         for var in global_parameters:
63             global_parameters[var] = (sum_parameters[var] / node_iter)
64     # test

```

```

61     net.load_state_dict(global_parameters, strict=True)
62     net.eval()
63     for data_set in nodes.test_loaders:
64         running_correct = 0
65         running_samples = 0
66         for data, label in data_set:
67             pred = net(data)
68             running_correct += pred.argmax(1).eq(label).sum().item()
69             running_samples += len(label)
70     print("\t" + 'accuracy: %.2f' % (running_correct / running_samples),
end="")

```

## 6. client训练函数

```

1  def local_upload(train_data_set, local_epoch, net, loss_fun, opt,
global_parameters, dev):
2      # 加载当前通信中最新全局参数
3      net.load_state_dict(global_parameters, strict=True)
4      # 设置迭代次数
5      net.train()
6      for epoch in range(local_epoch):
7          for data, label in train_data_set:
8              data, label = data.to(dev), label.to(dev)
9              # 模型上传入数据
10             predict = net(data)
11             loss = loss_fun(predict, label)
12             # 反向传播
13             loss.backward()
14             # 计算梯度，并更新梯度
15             opt.step()
16             # 将梯度归零，初始化梯度
17             opt.zero_grad()
18     # 返回当前Client基于自己的数据训练得到的新的模型参数
19     return net.state_dict()

```

## 7. 模型评估函数

```

1  def evaluate(net, global_parameters, testDataLoader, dev):
2      net.load_state_dict(global_parameters, strict=True)
3      running_correct = 0
4      running_samples = 0
5      net.eval()
6      # 载入测试集
7      for data, label in testDataLoader:
8          data, label = data.to(dev), label.to(dev)
9          pred = net(data)
10         running_correct += pred.argmax(1).eq(label).sum().item()
11         running_samples += len(label)
12     print("\t" + 'accuracy: %.2f' % (running_correct / running_samples),
end="")

```



## 8. 模型训练结果

因为设备原因，暂时无法训练出论文中的模型

Epoch: 2, Node Count: 35, Node ID: 39	accuracy: 0.7173376083374023
Epoch: 2, Node Count: 36, Node ID: 43	accuracy: 0.5102083086967468
Epoch: 2, Node Count: 37, Node ID: 41	accuracy: 0.6944901347160339
Epoch: 2, Node Count: 38, Node ID: 4	accuracy: 0.4871794879436493
Epoch: 2, Node Count: 39, Node ID: 47	accuracy: 0.46875
Epoch: 2, Node Count: 40, Node ID: 25	accuracy: 0.7411221265792847
Epoch: 2, Node Count: 41, Node ID: 39	accuracy: 0.6196151375770569
Epoch: 2, Node Count: 42, Node ID: 40	accuracy: 0.7840402126312256
Epoch: 2, Node Count: 43, Node ID: 16	accuracy: 0.7732007503509521
Epoch: 2, Node Count: 44, Node ID: 8	accuracy: 0.7512019276618958
Epoch: 2, Node Count: 45, Node ID: 27	accuracy: 0.6020998954772949
Epoch: 2, Node Count: 46, Node ID: 5	accuracy: 0.4984374940395355
Epoch: 2, Node Count: 47, Node ID: 40	accuracy: 0.7109375
Epoch: 2, Node Count: 48, Node ID: 49	accuracy: 0.7680289149284363

## 附录：关键函数记录

### torch.nn.Module.load\_state\_dict

**load\_state\_dict(state\_dict, strict=True)**

使用 state\_dict 反序列化模型参数字典。用来加载模型参数。将 state\_dict 中的 parameters 和 buffers 复制到此 module 及其子节点中。

概况：给模型对象加载训练好的模型参数，即加载模型参数

state\_dict (字典类型) – 一个包含参数和持续性缓冲的字典，往往是pytorch模型pth文件

strict (布尔类型, 可选) – 该参数用来指明是否需要强制严格匹配, 即:state\_dict中的关键字是否需要和该模块的state\_dict()方法返回的关键字强制严格匹配.默认值是True

### nn.utils.clip\_grad\_norm\_

**nn.utils.clip\_grad\_norm\_(parameters, max\_norm, norm\_type=2)**

这个函数是根据参数的范数来衡量的

Parameters:

parameters (Iterable[Variable]) – 一个基于变量的迭代器，会进行归一化（原文：an iterable of Variables that will have gradients normalized）

max\_norm (float or int) – 梯度的最大范数（原文：max norm of the gradients）

norm\_type(float or int) – 规定范数的类型，默认为L2（原文：type of the used p-norm. Can be'inf'for infinity norm）

Returns:参数的总体范数（作为单个向量来看）（原文：Total norm of the parameters (viewed as a single vector).）

### torch.nn.Embedding

```
1 torch.nn.Embedding(num_embeddings, embedding_dim, padding_idx=None,
    max_norm=None, norm_type=2.0, scale_grad_by_freq=False, sparse=False,
    _weight=None, device=None, dtype=None)
```

一个简单的查找表，用于存储固定字典和大小的嵌入。该模块通常用于存储词嵌入并使用索引检索它们。模块的输入是索引列表，输出是相应的词嵌入。

```

weights = {OrderedDict: 10} OrderedDict([('conv1.weight', tensor([[[[ 0.0494, 0.0171, -0.0424, -0.0269, 0.1247],\n
> 'conv1.weight' = {Tensor: (16, 3, 5, 5)} tensor([[[[ 0.0494, 0.0171, -0.0424, -0.0269, 0.1247],\n      [ 0.1498, 0.1479... View
> 'conv1.bias' = {Tensor: (16,)} tensor([[-0.0792, 0.0266, 0.0055, 0.1055, 0.0521, -0.0339, -0.0954, 0.0958,\n      0.1 ... View
> 'conv2.weight' = {Tensor: (32, 16, 5, 5)} tensor([[[[ 1.9384e-02, 6.0149e-02, 1.3496e-01, 2.9994e-03, -8.7224e-03],... View
> 'conv2.bias' = {Tensor: (32,)} tensor([ 1.1107e-02, 2.6046e-03, 9.7536e-02, 8.8637e-02, 2.4015e-02,\n      1.1596... View
> 'fc1.weight' = {Tensor: (120, 800)} tensor([[[ 0.1008, -0.0714, -0.0279, ..., -0.0402, -0.0823, -0.0188],\n      [ 0.0502, -... View
> 'fc1.bias' = {Tensor: (120,)} tensor([ 0.0327, 0.0043, -0.0101, 0.0872, -0.0259, -0.0517, -0.0444, -0.0744,\n      -0.00... View
> 'fc2.weight' = {Tensor: (84, 120)} tensor([[[ 0.1371, 0.0215, -0.1064, ..., -0.1323, 0.1125, -0.0415],\n      [ 0.0720, -0... View
> 'fc2.bias' = {Tensor: (84,)} tensor([[-0.0704, -0.0051, -0.0955, 0.1114, -0.0677, 0.0691, -0.0558, -0.0676,\n      -0.092... View
> 'fc3.weight' = {Tensor: (10, 84)} tensor([[[ 1.1749e-02, 1.6591e-01, -2.2738e-03, 2.4973e-02, 4.0403e-02,\n      -9... View
> 'fc3.bias' = {Tensor: (10,)} tensor([[-0.1472, -0.0609, 0.0361, 0.0724, 0.0070, -0.0630, 0.0898, 0.1310,\n      0.0817, 0.003
01 _len_ = {int} 10

```

```

CNNTarget(
  (conv1): Conv2d(3, 16, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
    ceil_mode=False)
  (conv2): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=800, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)

```